

Big Location Data:
Balancing Profits, Promise, and Perils
Thesis Proposal

Chris Riederer
Department of Computer Science
Columbia University
mani@cs.columbia.edu

February 9, 2017

Abstract

Ubiquitous, mobile computing in the form of smartphones has created data that lets us study human behavior like never before. In particular, data about human mobility has allowed us to understand the hows and whys of human movement. At the same time, these new collections of data can present societal risks, as we've now enabled mass surveillance, a loss of privacy, and algorithmic bias.

In this thesis proposal, I describe recent work that attempts to balance the scientific and engineering promises of location data with the potential risks. I will describe work I have completed relating location data to privacy, anonymity, economics, and algorithmic bias. I propose future research to be completed in the form of a thesis, advancing knowledge of location-based demographics and algorithmic bias.

Contents

1	Introduction	1
1.1	Outline	1
2	Background	1
2.1	Location Data	1
2.2	Privacy	2
2.3	Bias	3
2.4	Online Advertising	3
3	Location Data, Privacy, and Economics	3
3.1	Related Work	3
3.2	Completed Work	3
4	Location Data and Anonymity	4
4.1	Related Work	4
4.2	Completed Work	4
5	Location Data, Demographics, and Bias	5
6	Proposal Topic I	5
7	Proposal Topic II	6
8	Research plan	6

1 Introduction

TODO

1.1 Outline

I will begin with a background section which introduces the core concepts found in this proposal: location data, privacy, and bias. I proceed with three chapters detailing completed work. Each chapter contains a section summarizing relevant prior work.

Chapter 3 focuses on location data, privacy, and economics. We begin with work that seeks to understand user attitudes to their privacy and the economic value of their information. Specifically, it examines an alternative to the current practice of firms offering free services in exchange for full control over user data. The alternative model is one in which users control their data and make decisions about when to sell access to their info, and to whom.

Chapter 4 examines the possibility of anonymizing location data. Prior work has shown that users are highly unique in their location patterns, leaving them vulnerable to deanonymization (see 2.2). Here we take this a step further, showing not only that this vulnerability exists, but that users indeed can be linked to other datasets. Additionally, we provide a tool to users that aggregates and displays their location data along with the potential inferences made from it.

Chapter 5 shows the potential for location data to be part of systems that We gather a dataset of locations attached to demographic information from a popular image-sharing mobile application. This data allows us to study the differences in human mobility across different groups, and moreover, to show that demographics can be inferred using only location data. This raises questions about the sensitivity of location data, and about the potential for bias in systems that make decisions based on location data. We examine other methodologies for inferring demographics from social network data and discuss debiasing of algorithms.

Chapter 6

Chapter 7

I conclude with a plan for completing this work in Chapter 8.

2 Background

2.1 Location Data

What is location data? Most generally, location data is information relating people to places. Typically, this relation is the fact that a person was at a place. Adding time into the figure, the relation could be that a person was at a place at a particular time. However, location data could also include relations about the importance of a place in someones life, such as them living in a location, working at a location, or having spent a quantity of time in a location. Though location data does not need to be associated with user IDs, in this work we will consider that there is always attached some sort of user ID that uniquely identifies the user in the dataset, possibly de-personalized.

Location data can be described in two main ways: **geographically** or **semantically**. *Geographic* data can be described by a latitude-longitude data on the globe. *Semantic* location data refers to an identifier used within that dataset. This could have some information available to a

common user, e.g. “New York City”, or it could simply be an identifier, e.g. 7. Note that often these two may be combined or used together. A location such as “CEPSR Office 618, Columbia University” (the author’s office) indicates a very small, non-ambiguous location that can easily be mapped to geographic coordinates. Semantic location data can sometimes present a privacy problem, as an association with a place could indicate sensitive attributes, such as someone’s religion, political affiliation, health, or sexuality. In this work, I will typically assume location data is also tagged with temporal data, and I will use the terms location data and spatiotemporal data interchangeably.

To put this more formally, we can define a single data point p of location data to be:

$$p = \langle u, l \rangle$$

or, including time:

$$p = \langle u, l, t \rangle$$

where u uniquely identifies a user, l uniquely identifies a location, and t specifies a time. Note that l could be a latitude-longitude pair in the geographic case or an ID in the semantic case.

How is location data collected? Location data can be captured passively or actively. **Actively captured** location data is only recorded when the user takes some action. Note that this action does not need to inherently be “about” location data, for example, a user making a call from a cell phone or swiping a credit card is typically not consciously thinking about their location data. A record of their location is created as a by-product of their use of that technology. **Passively captured** is meant in a stronger way— the user’s location is captured without the user making any kind of action. This can occur through tracking apps. An example is MapMyRun¹, an app where users record their routes while running, in order to track distance and progress in meeting exercise goals. Although the user took an action to start recording their location, the location is recorded in the background with no user action from then on, and hence we call it “passive”. Another example is Google’s location history. Google records location data in the background of a users Android phone every few minutes. A map of everywhere a user (with an Android phone with location history turned on) is available at ².

What is location used for? TODO

2.2 Privacy

Privacy has been an important concept, brought to the forefront of public debate as surveillance of users has grown, both by governments and private companies.

ADD MORE STUFF

In this work, we will focus on two technical conceptions of privacy, *k-anonymity* and *differential privacy*.

k-anonymity

Differential privacy

¹<http://www.mapmyrun.com/>

²<https://www.google.co.in/maps/timeline>

2.3 Bias

2.4 Online Advertising

3 Location Data, Privacy, and Economics

The online economy is based primarily on advertising. The income of a firm roughly translates to (number of impressions) \times (dollars per impression). I am trying to keep this abstract and not saying that firms are always getting paid for impressions, as other models like paying per click or per sale or other action are quite common. Really the argument here is that firms make money based on how many people come to their site and how well they can target advertisements to those individuals. This gives firms an incentive to gather as much information about their users as possible so that they can better target ads to them.

This framework presents a challenge to privacy. User information is collected and gathered in one centralized place. There are multiple risks involved here: the firms themselves may use the information in ways the users disagree with, the firms may sell or be coerced to give their information to other firms or governments, or the firms may fall victim to cybersecurity attacks, leaking information to other sources. TODO: cite some stuff.

As ways to counter this, schemes have been proposed to encrypt user behavior and information, denying all access to a firm. However, this would deny firms the ability to make money, meaning no services would be provided for users and possibly a lower global utility be reached. Thus, schemes that ignore this economy however are unlikely to be adopted. Companies need to make money to function. Currently, users seem happy to provide their data in exchange for free services. A concern is that users do not have a good idea of their data and do not know how it is being used and to whom it is accessible.

Therefore it is important to gain an understanding of how users value their information, what they believe firms are doing with their data, and what users are comfortable with in terms of data use.

3.1 Related Work

3.2 Completed Work

How does one determine how a study participant values something as abstract as In our work, “Your Browsing Behavior for a Big Mac”, TODO(cite) we utilized

- Your Browsing Behavior for a Big Mac (maybe?)
- For Sale: Your Data. By: You (some portion?)
- Challenges of Keyword-Based Location Disclosure

“Your Browsing Behavior for a Big Mac” User privacy is extremely important. However, there does not exist a strong understanding of how users value their privacy.

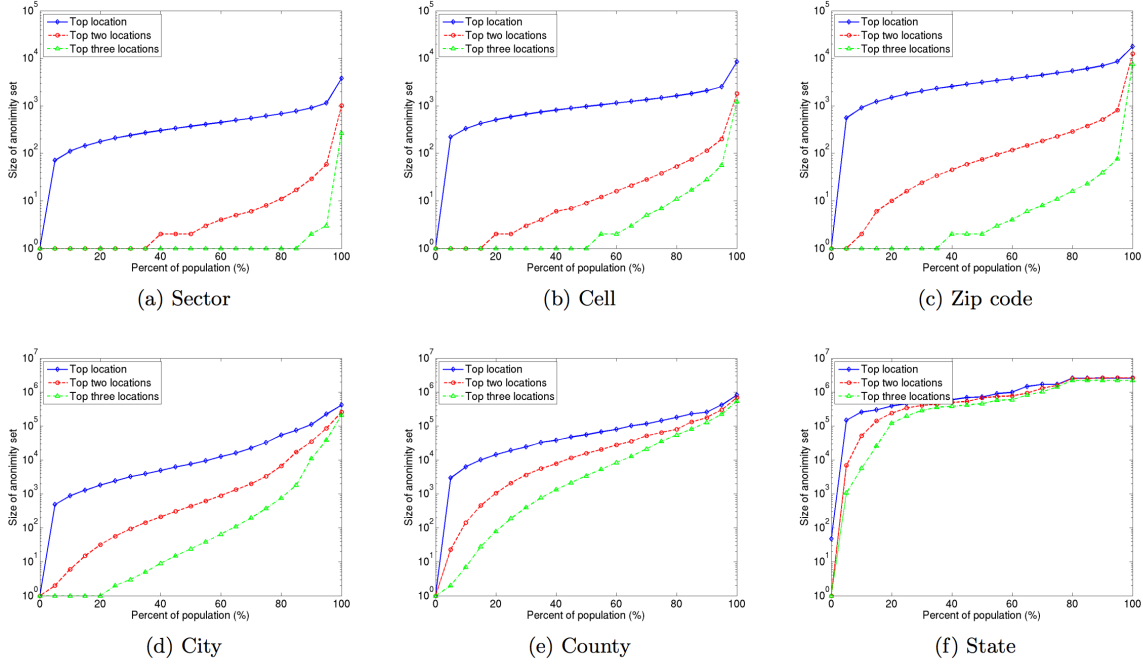


Figure 1: Figure from ?? depicting the size of anonymity sets for top n most visited location of users. Locations are varied in granularity, from cell sectors to US states.

4 Location Data and Anonymity

4.1 Related Work

Location data for individuals is highly unique and thus difficult to anonymize. The first large-scale study of the k -anonymity of location data was appropriately titled “Anonymization of Location Data Does Not Work” ???. The paper used data from cell phone call detail records (or CDR, see Chapter 2) for 25 million United States users over a 3 month period. The authors represents each user as simply their top n most visited locations, varying n from 1 to 3. Additionally, the authors varied the granularity of the locations, with the smallest as cell sector and the largest as state. Remarkably, using 3 locations at a cell level made half of all users completely unique, and 3 locations a sector level made 85% of all users unique. A figure detailing this result and results for other granularities and values of n is depicted in Figure 1. The authors went on to analyze the impact of geography (comparing different states and cities), mobility (distances between top locations), and social networks on anonymity.

The Montjoye nature report

4.2 Completed Work

I have investigated the anonymity of location data for users

Although prior work showed location to be highly *unique* and thus possibly *vulnerable* to de-anonymization, no data was actually de-anonymized in practice. Indeed, just because a data source

is highly unique does not mean it can be de-anonymized. For example, much of cryptography relies on creating highly unique but unpredictable sequences of numbers. To put it more concretely, imagine that each individual had a die with 1000 sides, and each side represented a location. If, quite hypothetically, humans decided where to go next by rolling this die, their movements would look very unique. However, since the movements are random and unpredictable, my movements from different time periods will be indistinguishable from those of a different individual.

TODO: put some math here?

Another possible break in the argument that uniqueness implies vulnerability is the important factor of sampling. The datasets dealt with here (phone records, social media posts) are all *actively* collected: each data point exists if and only if the user has taken an action. Intuitively, the location data from different sampling data sources should look very different. An individual may be more likely to make phone calls in quiet places, like the home or office, and take geotagged location photos in popular tourist destinations or restaurants.

TODO: put some math here?

In “Linking Users Across Domains with Location Data”, published at WWW in 2016, we tackled this problem, linking users across two entirely different datasets.

To conduct this work, we obtained three datasets. This in itself was a significant challenge, as each dataset needed to contain individuals with identities linked across two different data sources.

- Cell phone-Credit Card
- Instagram-Twitter
- Foursquare-Twitter

and

FindYou

5 Location Data, Demographics, and Bias

I Dont Have a Photograph But You Can Have My Footprints Under submission work on demographic labeling Current work on bias!

6 Proposal Topic I

As described in Chapter 5, an important challenge facing the computer science community is algorithmic bias. In recent years, an emerging body of work has focused on different mitigating techniques, such as automated discovery of bias, “de-biasing” existing algorithms, or theoretical analyses of different types of bias. De-biasing techniques are sure to incur a cost: the objective function of the algorithm is no longer as straightforward, and organizationally new infrastructure needs to be put into place for something that could hurt revenue. Understanding the key trade-offs between revenue and uncertain risk will be important to insure real-world adoption. Although there have been some good initial insights, the community has lacked strong data-driven analysis on this trade off.

I propose to fill this gap by applying proposed techniques to real-world problems through the use of an innovative dataset. Namely, I will look at the real-world problems of recommendation systems within a large social network. I will examine the trade off between recommendation accuracy, bias, and revenue.

Over the course of several months I have gathered photo metadata from the popular image-sharing application Instagram. I have run these photos through a program that recognizes faces within each image, tagging it with age, gender, and ethnicity. This will create the largest publicly available dataset that I know of connecting human mobility to demographics.

Machine learning systems utilize location in making recommendations. However, location can be highly correlated with potentially sensitive traits, such as ethnicity. I plan to look at

The project will emerge in several stages.

1. Collection of instagram data (completed).
2. Labeling of instagram data with Face++ API (completed).
3. Initial analysis and descriptive statistics of dataset (in progress).
4. Full problem specification: algorithms, inputs, and objectives.
5. Apply de-biasing to algorithms and analyze impacts.
6. Create recommendations for algorithm designers.

7 Proposal Topic II

8 Research plan